Lecture 3: Planning by Dynamic Programming

# Lecture 3: Planning by Dynamic Programming

David Silver

#### Outline

- 1 Introduction
- 2 Policy Evaluation
- 3 Policy Iteration
- 4 Value Iteration
- 5 Extensions to Dynamic Programming
- 6 Contraction Mapping

# What is Dynamic Programming?

Dynamic sequential or temporal component to the problem Programming optimising a "program", i.e. a policy

- c.f. linear programming
- A method for solving complex problems
- By breaking them down into subproblems
  - Solve the subproblems
  - Combine solutions to subproblems

## Requirements for Dynamic Programming

Dynamic Programming is a very general solution method for problems which have two properties:

- Optimal substructure
  - Principle of optimality applies
  - Optimal solution can be decomposed into subproblems
- Overlapping subproblems
  - Subproblems recur many times
  - Solutions can be cached and reused
- Markov decision processes satisfy both properties
  - Bellman equation gives recursive decomposition
  - Value function stores and reuses solutions

# Planning by Dynamic Programming

- Dynamic programming assumes full knowledge of the MDP
- It is used for planning in an MDP
- For prediction:
  - Input: MDP  $\langle S, A, P, R, \gamma \rangle$  and policy  $\pi$
  - or: MRP  $\langle \mathcal{S}, \mathcal{P}^{\pi}, \mathcal{R}^{\pi}, \gamma \rangle$
  - Output: value function  $v_{\pi}$
- Or for control:
  - Input: MDP  $\langle \mathcal{S}, \mathcal{A}, \mathcal{P}, \mathcal{R}, \gamma \rangle$
  - lacksquare Output: optimal value function  $v_*$
  - lacksquare and: optimal policy  $\pi_*$

# Other Applications of Dynamic Programming

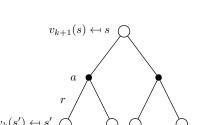
Dynamic programming is used to solve many other problems, e.g.

- Scheduling algorithms
- String algorithms (e.g. sequence alignment)
- Graph algorithms (e.g. shortest path algorithms)
- Graphical models (e.g. Viterbi algorithm)
- Bioinformatics (e.g. lattice models)

# Iterative Policy Evaluation

- lacktriangle Problem: evaluate a given policy  $\pi$
- Solution: iterative application of Bellman expectation backup
- $\longrightarrow v_2 \rightarrow ... \rightarrow v_\pi$
- Using synchronous backups,
  - At each iteration k+1
- For all states  $s \in \mathcal{S}$ 
  - Update  $v_{k+1}(s)$  from  $v_k(s')$
  - $\blacksquare$  where s' is a successor state of s
- We will discuss asynchronous backups later
- $lue{}$  Convergence to  $v_{\pi}$  will be proven at the end of the lecture

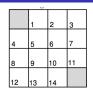
# Iterative Policy Evaluation (2)



$$\begin{aligned} v_{k+1}(s) &= \sum_{a \in \mathcal{A}} \pi(a|s) \left( \mathcal{R}_s^a + \gamma \sum_{s' \in \mathcal{S}} \mathcal{P}_{ss'}^a v_k(s') \right) \\ \mathbf{v}^{k+1} &= \mathcal{R}^{\pi} + \gamma \mathcal{P}^{\pi} \mathbf{v}^k \end{aligned}$$

# Evaluating a Random Policy in the Small Gridworld





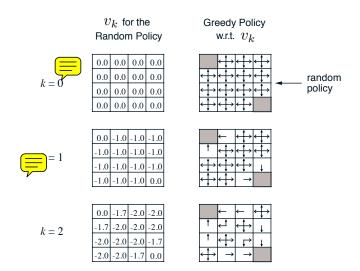
r = -1 on all transitions

- Undiscounted episodic MDP  $(\gamma = 1)$
- Nonterminal states 1, ..., 14
- One terminal state (shown twice as shaded squares)
- Actions leading out of the grid leave state unchanged
- $\blacksquare$  Reward is -1 until the terminal state is reached
- Agent follows uniform random policy

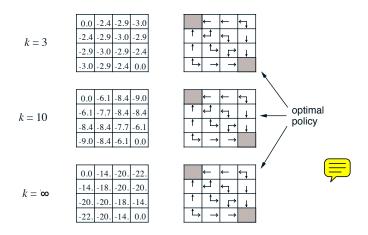


$$\pi(n|\cdot) = \pi(e|\cdot) = \pi(s|\cdot) = \pi(w|\cdot) = 0.25$$

### Iterative Policy Evaluation in Small Gridworld



# Iterative Policy Evaluation in Small Gridworld (2)



## How to Improve a Policy

- Given a policy  $\pi$ 
  - **Evaluate** the policy  $\pi$



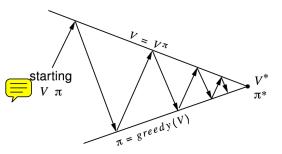
$$v_{\pi}(s) = \mathbb{E}[R_{t+1} + \gamma R_{t+2} + ... | S_t = s]$$

**Improve** the policy by acting greedily with respect to  $v_{\pi}$ 

$$\pi' = \mathsf{greedy}(v_\pi)$$

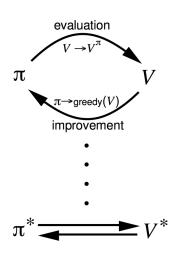
- In Small Gridworld improved policy was optimal,  $\pi' = \pi^*$
- In general, need more iterations of improvement / evaluation
- But this process of policy iteration always converges to  $\pi*$

#### Policy Iteration



Policy evaluation Estimate  $v_{\pi}$  Iterative policy evaluation

Policy improvement Generate  $\pi' \geq \pi$  Greedy policy improvement



Policy Iteration

Example: Jack's Car Rental

#### Jack's Car Rental

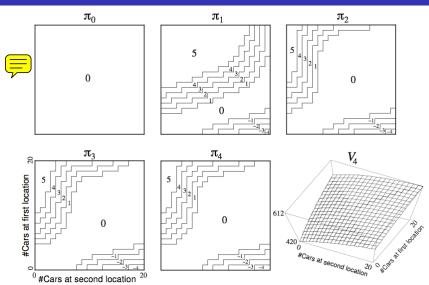


- States: Two locations, maximum of 20 cars at each
- Actions: Move up to 5 cars between locations overnight
- Reward: \$10 for each car rented (must be available)
- Transitions: Cars returned and requested randomly
  - Poisson distribution, *n* returns/requests with prob  $\frac{\lambda^n}{n!}e^{-\lambda}$
  - 1st location: average requests = 3, average returns = 3
  - 2nd location: average requests = 4, average returns = 2

Policy Iteration

Example: Jack's Car Rental

# Policy Iteration in Jack's Car Rental



# Policy Improvement

- Consider a deterministic policy,  $a = \pi(s)$
- We can *improve* the policy by acting greedily

$$\frac{\pi'(s) = \operatorname{argmax} q_{\pi}(s, a)}{a \in \mathcal{A}}$$

■ This improves the value from any state *s* over one step,

$$q_\pi(s,\pi'(s)) = \max_{a\in\mathcal{A}} q_\pi(s,a) \geq q_\pi(s,\pi(s)) = v_\pi(s)$$



lacksquare It therefore improves the value function,  $v_{\pi'}(s) \geq v_{\pi}(s)$ 

$$\begin{aligned} v_{\pi}(s) &\leq q_{\pi}(s, \pi'(s)) = \mathbb{E}_{\pi'} \left[ R_{t+1} + \gamma v_{\pi}(S_{t+1}) \mid S_{t} = s \right] \\ &\leq \mathbb{E}_{\pi'} \left[ R_{t+1} + \gamma q_{\pi}(S_{t+1}, \pi'(S_{t+1})) \mid S_{t} = s \right] \\ &\leq \mathbb{E}_{\pi'} \left[ R_{t+1} + \gamma R_{t+2} + \gamma^{2} q_{\pi}(S_{t+2}, \pi'(S_{t+2})) \mid S_{t} = s \right] \\ &\leq \mathbb{E}_{\pi'} \left[ R_{t+1} + \gamma R_{t+2} + \dots \mid S_{t} = s \right] = v_{\pi'}(s) \end{aligned}$$

# Policy Improvement (2)

If improvements stop,

$$q_\pi(s,\pi'(s)) = \max_{a\in\mathcal{A}} q_\pi(s,a) = q_\pi(s,\pi(s)) = v_\pi(s)$$

Then the Bellman optimality equation has been satisfied

$$v_{\pi}(s) = \max_{a \in \mathcal{A}} q_{\pi}(s, a)$$

- lacksquare Therefore  $v_\pi(s)=v_*(s)$  for all  $s\in\mathcal{S}$
- lacksquare so  $\pi$  is an optimal policy

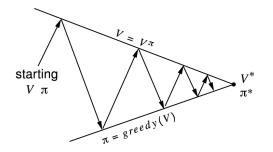


# Modified Policy Iteration

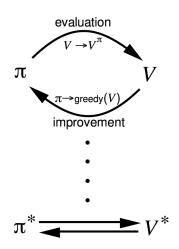


- Does policy evaluation need to converge to  $v_{\pi}$ ?
- Or should we introduce a stopping condition
  - lacksquare e.g.  $\epsilon$ -convergence of value function
- Or simply stop after k iterations of iterative policy evaluation?
- For example, in the small gridworld k = 3 was sufficient to achieve optimal policy
- Why not update policy every iteration? i.e. stop after k = 1
  - This is equivalent to *value iteration* (next section)

## Generalised Policy Iteration



Policy evaluation Estimate  $v_{\pi}$ Any policy evaluation algorithm Policy improvement Generate  $\pi' \geq \pi$ Any policy improvement algorithm



# Principle of Optimality

Any optimal policy can be subdivided into two components:

- An optimal first action A<sub>\*</sub>
- $lue{}$  Followed by an optimal policy from successor state S'

#### Theorem (Principle of Optimality)

A policy  $\pi(a|s)$  achieves the optimal value from state s,  $v_{\pi}(s) = v_{*}(s)$ , if and only if

- For any state s' reachable from s
- lacktriangledown  $\pi$  achieves the optimal value from state s',  $v_\pi(s')=v_*(s')$

#### Deterministic Value Iteration

- If we know the solution to subproblems  $v_*(s')$
- Then solution  $v_*(s)$  can be found by one-step lookahead

$$v_*(s) \leftarrow \max_{a \in \mathcal{A}} \mathcal{R}_s^a + \gamma \sum_{s' \in \mathcal{S}} \mathcal{P}_{ss'}^a v_*(s')$$

- The idea of value iteration is to apply these updates iteratively
- Intuition: start with final rewards and work backwards
- Still works with loopy, stochastic MDPs

# Example: Shortest Path

g		

0	0	0	0
0	0	0	0
0	0	0	0
0	0	0	0

0	-1	-1	-1
-1	-1	-1	-1
-1	-1	-1	-1
-1	-1	7	-1

-1	-2	-2
-2	-2	-2
-2	-2	-2
-2	-2	-2
	-2	-2 -2 -2 -2

Problem

V<sub>1</sub>

V<sub>2</sub>

 $V_3$ 

0	-1	-2	-3
-1	-2	-3	-3
-2	-3	-3	-3
-3	-3	-3	-3
$V_4$			

0	-1	-2	-3	
-1	-2	-3	-4	
-2	-3	-4	-4	
-3	-4	-4	-4	
V <sub>5</sub>				

-1	-2	-3
-2	-3	-4
-3	-4	-5
-4	-5	-5
	-2	-2 -3 -3 -4

0	-1	-2	-3
-1	-2	-3	-4
-2	-3	-4	-5
-3	-4	-5	-6

#### Value Iteration



■ Problem: find optimal policy  $\pi$ 

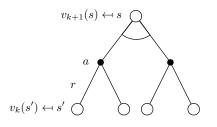


- Solution: iterative application of Beliman optimality backup
- $\mathbf{v}_1 \rightarrow \mathbf{v}_2 \rightarrow ... \rightarrow \mathbf{v}_*$
- Using synchronous backups
  - At each iteration k+1
  - lacksquare For all states  $s \in \mathcal{S}$
  - Update  $v_{k+1}(s)$  from  $v_k(s')$
- Convergence to  $v_*$  will be proven later
- Unlike policy iteration, there is no explicit policy
- Intermediate value functions may not correspond to any policy

└─Value Iteration

└Value Iteration in MDPs

# Value Iteration (2)



$$\begin{aligned} v_{k+1}(s) &= \max_{a \in \mathcal{A}} \left( \mathcal{R}_s^a + \gamma \sum_{s' \in \mathcal{S}} \mathcal{P}_{ss'}^a v_k(s') \right) \\ \mathbf{v}_{k+1} &= \max_{a \in \mathcal{A}} \mathcal{R}^a + \gamma \mathcal{P}^a \mathbf{v}_k \end{aligned}$$

└Value Iteration in MDPs

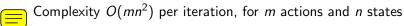
#### Example of Value Iteration in Practice

 $http://www.cs.ubc.ca/{\sim}poole/demos/mdp/vi.html$ 

# Synchronous Dynamic Programming Algorithms

	Problem	Bellman Equation	Algorithm
	Duadiation	Bellman Expectation Equation	Iterative
Prediction	Beillian Expectation Equation	Policy Evaluation	
	Control	Bellman Expectation Equation + Greedy Policy Improvement	Policy Iteration
7	Control	Bellman Optimality Equation	Value Iteration

■ Algorithms are based on state-value function  $v_{\pi}(s)$  or  $v_{*}(s)$ 



- Could also apply to action-value function  $q_{\pi}(s,a)$  or  $q_*(s,a)$
- Complexity  $O(m^2n^2)$  per iteration

# Asynchronous Dynamic Programming



- DP methods described so far used *synchronous* backups
- i.e. all states are backed up in parallel
- Asynchronous DP backs up states individually, in any order
- For each selected state, apply the appropriate backup
- Can significantly <u>reduce computation</u>
- Guaranteed to converge if all states continue to be selected

# Asynchronous Dynamic Programming

Three simple ideas for <u>asynchronous dynamic programming</u>:

- In-place dynamic programming
- Prioritised sweeping



■ Real-time dynamic programming

# In-Place Dynamic Programming

Synchronous value iteration stores two copies of value function for all s in  $\mathcal S$ 

$$v_{new}(s) \leftarrow \max_{a \in \mathcal{A}} \left( \mathcal{R}_{s}^{a} + \gamma \sum_{s' \in \mathcal{S}} \mathcal{P}_{ss'}^{a} v_{old}(s') \right)$$

$$V_{old} \leftarrow V_{new}$$

In-place value iteration only stores one copy of value function

$$\frac{\text{for all } s \text{ in } \mathcal{S}}{\mathsf{v}(s)} \leftarrow \max_{s \in \mathcal{A}} \left( \mathcal{R}_s^s + \gamma \sum_{s' \in \mathcal{S}} \mathcal{P}_{ss'}^a \mathsf{v}(s') \right)$$

# Prioritised Sweeping

■ Use magnitude of Bellman error to guide state selection, e.g.

$$\left| \max_{\mathbf{a} \in \mathcal{A}} \left( \mathcal{R}_{s}^{\mathbf{a}} + \gamma \sum_{s' \in \mathcal{S}} \mathcal{P}_{ss'}^{\mathbf{a}} v(s') \right) - v(s) \right| \qquad \boxed{=}$$

- Backup the state with the largest remaining Bellman error
- Update Bellman error of affected states after each backup
- Requires knowledge of reverse dynamics (predecessor states)
- Can be implemented efficiently by maintaining a priority queue

#### Asynchronous Dynamic Programming

# Real-Time Dynamic Programming



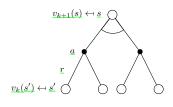
- Idea: only states that are relevant to agent
- Use agent's experience to guide the selection of states
- After each time-step  $S_t$ ,  $A_t$ ,  $R_{t+1}$
- $\blacksquare$  Backup the state  $S_t$

$$v(S_t) \leftarrow \max_{a \in \mathcal{A}} \left( \mathcal{R}_{S_t}^a + \gamma \sum_{s' \in \mathcal{S}} \mathcal{P}_{S_t s'}^a v(s') \right)$$

#### Full-Width Backups

- DP uses full-width backups
- For each backup (sync or async)
  - Every successor state and action is considered
  - Using knowledge of the MDP transitions and reward function
- DP is effective for medium-sized problems (millions of states)
- For large problems DP suffers Bellman's curse of dimensionality
  - Number of states n = |S| grows exponentially with number of state variables
- Even one backup can be too expensive



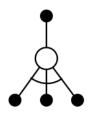


# Sample Backups



- In subsequent lectures we will consider <u>sample backups</u>
- Using sample rewards and sample transitions  $\langle S, A, R, S' \rangle$
- $lue{}$  Instead of reward function  ${\cal R}$  and transition dynamics  ${\cal P}$
- Advantages:
  - Model-free: no advance knowledge of MDP required
  - Breaks the curse of dimensionality through sampling
  - Cost of backup is constant, independent of n = |S|





# Approximate Dynamic Programming

- Approximate the value function
- Using a function approximator  $\hat{v}(s, \mathbf{w})$
- Apply dynamic programming to  $\hat{v}(\cdot, \mathbf{w})$
- $\blacksquare$  e.g. Fitted Value Iteration repeats at each iteration k,
  - lacksquare Sample states  $ilde{\mathcal{S}} \subseteq \mathcal{S}$
  - For each state  $s \in \tilde{\mathcal{S}}$ , estimate target value using Bellman optimality equation,

$$\tilde{v}_k(s) = \max_{a \in \mathcal{A}} \left( \mathcal{R}_s^a + \gamma \sum_{s' \in \mathcal{S}} \mathcal{P}_{ss'}^a \hat{v}(s', \mathbf{w_k}) \right)$$

■ Train next value function  $\hat{v}(\cdot, \mathbf{w_{k+1}})$  using targets  $\{\langle s, \tilde{v}_k(s) \rangle\}$ 

#### Some Technical Questions

- How do we know that value iteration converges to  $v_*$ ?
- Or that iterative policy evaluation converges to  $v_{\pi}$ ?
- And therefore that policy iteration converges to  $v_*$ ?
- Is the solution unique?
- How fast do these algorithms converge?
- These questions are resolved by contraction mapping theorem

### Value Function Space

- lacksquare Consider the vector space  ${\mathcal V}$  over value functions
- There are |S| dimensions
- **Each** point in this space fully specifies a value function v(s)
- What does a Bellman backup do to points in this space?
- We will show that it brings value functions *closer*
- And therefore the backups must converge on a unique solution

#### Value Function ∞-Norm

- We will measure distance between state-value functions u and v by the  $\infty$ -norm
- i.e. the largest difference between state values,

$$||u-v||_{\infty} = \max_{s \in \mathcal{S}} |u(s)-v(s)|$$

## Bellman Expectation Backup is a Contraction

■ Define the Bellman expectation backup operator  $T^{\pi}$ ,

$$T^{\pi}(\mathbf{v}) = \mathcal{R}^{\pi} + \gamma \mathcal{P}^{\pi} \mathbf{v}$$

■ This operator is a  $\gamma$ -contraction, i.e. it makes value functions closer by at least  $\gamma$ ,

$$||T^{\pi}(u) - T^{\pi}(v)||_{\infty} = ||(\mathcal{R}^{\pi} + \gamma \mathcal{P}^{\pi} u) - (\mathcal{R}^{\pi} + \gamma \mathcal{P}^{\pi} v)||_{\infty}$$

$$= ||\gamma \mathcal{P}^{\pi}(u - v)||_{\infty}$$

$$\leq ||\gamma \mathcal{P}^{\pi}||u - v||_{\infty}||_{\infty}$$

$$\leq \gamma ||u - v||_{\infty}$$

### Contraction Mapping Theorem

#### Theorem (Contraction Mapping Theorem)

For any metric space V that is complete (i.e. closed) under an operator T(v), where T is a  $\gamma$ -contraction,

- T converges to a unique fixed point
- lacktriangle At a linear convergence rate of  $\gamma$

# Convergence of Iter. Policy Evaluation and Policy Iteration

- The Bellman expectation operator  $T^{\pi}$  has a unique fixed point
- $v_{\pi}$  is a fixed point of  $T^{\pi}$  (by Bellman expectation equation)
- By contraction mapping theorem
- Iterative policy evaluation converges on  $v_{\pi}$
- Policy iteration converges on v<sub>\*</sub>

### Bellman Optimality Backup is a Contraction

■ Define the Bellman optimality backup operator T\*,

$$T^*(v) = \max_{a \in \mathcal{A}} \mathcal{R}^a + \gamma \mathcal{P}^a v$$

■ This operator is a  $\gamma$ -contraction, i.e. it makes value functions closer by at least  $\gamma$  (similar to previous proof)

$$||T^*(u) - T^*(v)||_{\infty} \le \gamma ||u - v||_{\infty}$$

## Convergence of Value Iteration

- The Bellman optimality operator *T\** has a unique fixed point
- $lackbox{v}_*$  is a fixed point of  $\mathcal{T}^*$  (by Bellman optimality equation)
- By contraction mapping theorem
- Value iteration converges on  $v_*$